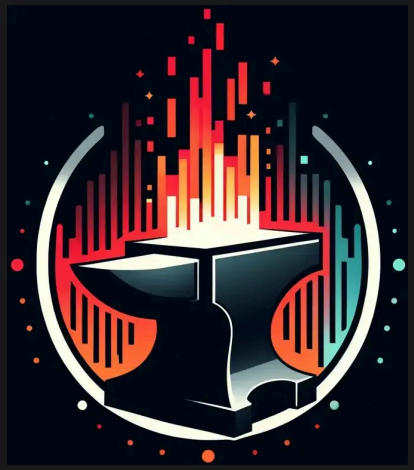

SMITH

Unleashing Neural Network Performance Across All Hardware

THE CHALLENGE

- Diverse hardware ecosystems
- Dominance of specific APIs like CUDA
- Customers forced to rely on NVIDIA GPUs

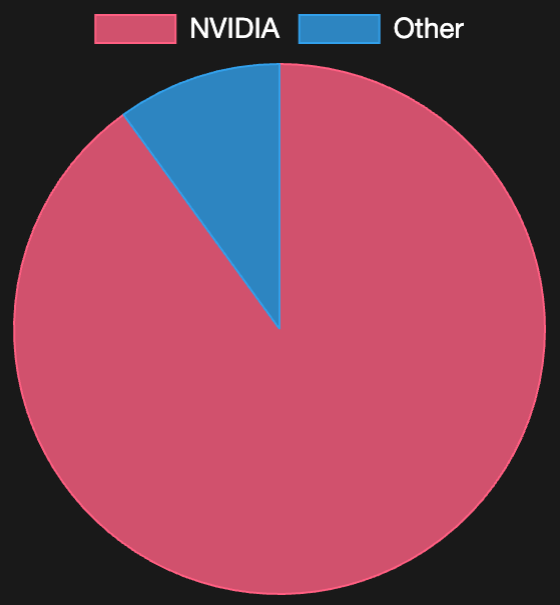

CURRENT STATE OF THE ART

- TensorFlow and PyTorch dominate
- Rely on manually optimized kernels
- High performance confined to specific platforms

	Nvidia	Amd	Intel	Huawei	...	Apple
Kernel 1	✓	✓		✓		
Kernel 2	✓	✓		✓		✓
Kernel 3	✓		✓	✓		✓
Kernel 4	✓		✓			✓
...						
Kernel N		✓	✓	✓		✓

INTRODUCING SMITH

- A single, open-source language
- Compiles to CPU and GPU backends
- Ensures optimal performance on all devices
- Targets clients, servers, IoT devices, robots

ROADMAP AND MILESTONES

- Year 1: WebAssembly for CPU support
- Year 2: WebGPU for GPU support
- Future targets: XLA, LLVM, etc

FUNDING OBJECTIVES

- Funding goal: $10 million
- Enable full-time development on Smith
- Facilitate open source ecosystem

JOIN SMITH

Invest in the future of neural networks. Help diversify hardware utilization in AI. Strengthen and expand market reach.